



DEVELOPER HANDOFF

Ancient Covenant Scrolls / Load Suite

Architecture · build state · how to move forward

git tip 720df39 (post-cleanup) · last verified state v17e0 (593d410)

1. 60-second brief
2. Repository structure
3. The five PWAs
4. Load — architecture
5. Image Prompt — architecture
6. Runtime patterns
7. What is shipped (v17dn → v17e9)
8. What is verified vs not
9. Open roadmap (37 items)
10. How to add a feature
11. How to add a provider
12. Surgical revert procedure
13. Cache discipline + deploy flow
14. Backup branches
15. Source-of-truth doc map
16. Gotchas / capability gaps
17. Recommended next steps

1. 60-second brief

This is a self-contained **multi-PWA suite** hosted on GitHub Pages, built to run offline-first on iPad Safari. Five active apps under one repo:

- **Load** (/load) — PWA host, code editor, image studio, video editor, exporter
- **Attain** (/attain) — dyslexia-friendly universal book reader + study activities
- **Attain Jr** (/attain-jr) — kid variant; COPPA-friendly; not yet user-verified
- **Study** (/study) — ACR study companion
- **ACR Reader** (root /) — Ancient Covenant Record reader
- **Great Eraser** (/GreatE) — companion bonus app

Active focus: **Image Prompt** (/load/image-prompt) — multi-provider AI image studio inside Load, Glam-AI-class chat experience with mask-painter, character lock, output verification, and image-to-video on the roadmap.

Hard rule, full stop: the user is dyslexic and works on iPad. Tone for any future helper: plain, short steps, no walls of text, never overpromise. Cache strings always go forward. Surgical revert only. No iPad-side WASM without a memory test.

2. Repository structure

```
ancient-covenant-scrolls/
├─ index.html           ← root page = ACR Reader (live at /)
├─ load/                ← Load PWA (with image-prompt sub-PWA)
│   ├─ index.html, load.js, load.css, sw.js, manifest.json
│   ├─ lib-pdf*.min.js, lib-epub.min.js, lib-jszip.min.js, lib-hebrew-pron.js
│   ├─ fonts/ (Atkinson + OpenDyslexic woff2)
│   ├─ image-prompt/    ← active feature (the AI image studio)
│   │   ├─ index.html ← single-file PWA, ~16k lines including JS
│   │   ├─ sw.js, manifest.json, icon.png, fonts/
│   │   └─ lib-* (vendored when needed)
│   ├─ book-video/      ← Book-to-Video module (specced, not wired)
│   └─ tools/           ← build helpers
├─ attain/              ← Attain PWA
├─ attain-jr/           ← Attain Jr PWA
├─ study/               ← Study PWA
├─ GreatE/              ← Great Eraser companion
├─ inbox/               ← user uploads + handoff zips + manuals
├─ PLAN_LOAD_AI.md      ← merged build plan v10 (single source of truth)
├─ PLAN_IMAGE_PROMPT.md, PLAN_IMAGE_PROMPT_v3.md, PLAN_BOOK_TO_VIDEO.md
├─ HANDOFF.md           ← long-running cross-session task list
├─ VERIFIED_LOG.md      ← per-version verification ledger
├─ CLAUDE.md            ← rules for AI / human developers
├─ SESSION_NOTES_*.md   ← per-day session journals
├─ designs/, docs/, data/, load-ios/
└─ *.zip / *.html       ← packaged downloadable artifacts
```

3. The five PWAs

App	Path	Cache namespace	Purpose
ACR Reader	/	acr-vNN	Bible-style reader (root index.html)
Load	/load	load-vXXY	PWA host + AI image + code editor

Image Prompt	<code>/load/image-prompt</code>	<code>image-prompt-vNN</code>	Sub-PWA — Glam-style AI image studio
Attain	<code>/attain</code>	<code>attain-vNN</code>	Dyslexia-friendly reader + study
Attain Jr	<code>/attain-jr</code>	<code>attain-jr-vNN</code>	Kid variant (in progress)
Study	<code>/study</code>	<code>acr-study-vNN</code>	ACR-aware study companion
Great Eraser	<code>/GreatE</code>	n/a	Companion bonus PWA

Each app has its own service worker and cache namespace. They never collide. Bumping `load-v17e9` doesn't affect Attain.

4. Load — architecture

`load/load.js` is one ~25k-line vanilla ES5 file (no transpiler, no build step). Hand-edited; runs as-is in iPad Safari. Top-level structure:

- **State:** settings, currentApp, currentBlobUrl, IndexedDB handle
- **Screen system:** home / library / import / settings / viewer / editor / note-editor — each is a section with class `screen`, toggled via `show()` / `hide()`
- **Library:** `apps[]` array kept in IndexedDB; `putApp()` / `getApps()` async helpers
- **Viewer:** blob-URL iframe via `openApp()` at `load.js:17349`
- **Code editor:** textarea-based editor with light/dark/cream themes; reaches into the active app's HTML
- **Image Prompt:** separate sub-PWA loaded via fullscreen iframe overlay from Load's Workspace tile (🎨)
- **Workspace tiles:** the toolbox grid in Library — Code Editor, Prose Editor, Cover Designer, Video Editor, Asset Doctor, File Tree, Manifest, Service Worker, AI Copilot, exports (EPUB / PDF / Standalone HTML / PWA Zip / App Package), Image Prompt
- **AI Copilot:** slash-command helper (`/explain` , `/fix` , `/optimize` , `/plan` , `/analyze`) — context-aware over the active app file

Import pipeline (zip path)

```
file picker
  → JSZip.loadAsync(file)
  → find index.html (root or first match)
  → for each non-index file: { type: 'text', body }
    | { type: 'binary', mime, body: base64 }
  → injectBundleShim(indexHtml, bundle, baseDir)
    runtime fetch / XHR / Image / script interceptor
    neuters service-worker registrations
  → app = { id, name, kind: 'zip', html, bundleIndex, bundleSamples }
  → putApp(app) → IndexedDB
```

Viewer pipeline

```
openApp(app):
  htmlBlob = new Blob([app.html], { type: 'text/html; charset=utf-8' })
  currentBlobUrl = URL.createObjectURL(htmlBlob)
  iframe.src = currentBlobUrl
  show('viewer-screen')
```

The injected shim makes the iframe app think it's running normally — fetches resolve from the embedded bundle dict.

5. Image Prompt — architecture

load/image-prompt/index.html — single-file PWA, all JS inline (no separate .js). The active build target.

The 6-pipeline spec (from PLAN_IMAGE_PROMPT_v3.md)

- **A. Vision Analysis** — vision LLM extracts structured JSON profile from uploaded image (status: was shipped v17dh, rolled back; needs re-ship)
- **B. Prompt Builder** — strict / moderate / loose modes (status: strict via Scene Lock; 3-mode UI open)
- **C. Image Editing** — img2img / inpaint / mask / object-add / object-remove / bg-change (status: Manual Mask painter shipped + verified; ADD_OBJECT auto-placement shipped)
- **D. Image Generation** — text-to-image multi-provider (status: 17 providers in chain)

- **E. Output Verification** — vision diff post-gen, auto-retry below threshold (status: open)
- **F. Fallback** — export-prompt for ComfyUI / A1111 / Fooocus / Runway / Kling / Pika / Luma / Wan when no capable provider (status: shipped for image; video fallback open)

Key inline functions (all in `load/image-prompt/index.html`)

Function	~ Line	Role
<code>classifyImageTask(prompt, hasImage)</code>	1845	Maps user prompt → task type (generate / add_object / remove_object / change_background / inpaint / face_restore / upscale / improve / modify_character)
<code>filterImageProvidersForTask(providers, task)</code>	1863	Strips taskOnly + ineligible providers per task
<code>runImageTask(opts)</code>	~1880	Single facade — accepts {prompt, sourceImage, mask, characterProfile, taskType}; routes
<code>imageGenWithFallback(prompt, imgData, opts)</code>	~2730	Provider chain executor with capability gate, withTimeout, skipped-reasons
<code>hfImageCall(model, prompt, imgData, mask)</code>	~2225	HF Inference caller with HF_TEXT_ONLY_MODELS pre-block
<code>callPuter / callGemini / callOR / callCerebras / callAnthropic / callHFChat</code>	1700-1830	Per-provider chat call, response normalized at source
<code>autoPlaceForObject(object)</code>	~1900	Returns {region, box[x y w h], hint} for ADD_OBJECT
<code>imageSimilarity(srcUrl, resUrl)</code>	~1930	8x8 luma diff for ADD_OBJECT no-change rejection
<code>buildSDEExportPrompt(prompt, lastSDPrompt)</code>	~1750	ComfyUI / A1111 / Fooocus export
<code>maskKey(k)</code>	~570	Renders <code>hf_****abcd</code> for Test Keys

<code>runKeyDiagnostic()</code>	~620	Test Keys driver (parallel <code>checkOneKey</code> calls)
<code>recordReceipt({ task, providerId, ... })</code>	~750	Output Receipt audit log
<code>openMask / doneMask</code>	~1075	Manual Mask painter (Canvas-based Editor Tool)

State globals

<code>C</code>	Settings + provider keys (loaded from <code>localStorage</code>)
<code>img</code>	Currently-attached image as <code>{dataUrl, base64, mime, name}</code>
<code>currentMask</code>	Painted mask as <code>{dataUrl, base64, mime, paintedPixels}</code> or null
<code>CHAR</code>	Character profile getters/setters into <code>localStorage</code> <code>ps_char_*</code>
<code>history</code>	Chat conversation history
<code>lastSDPrompt</code>	Last <code>SD_PROMPT</code> for multi-turn preservation context
<code>sceneLocked</code>	Scene Lock toggle state
<code>generating</code>	UI lock during in-flight gen
<code>PROVIDER_HEALTH</code>	Circuit-breaker: fail/success counts per provider

6. Runtime patterns

6.1 Capability registry — the lock

Every chat provider in `MODEL_CAPABILITIES` ; every image provider entry in `IMAGE_PROVIDERS[]` carries flags:

```
// chat (load/image-prompt/index.html, line ~575)
puter:    { text_only: false, vision_input: true, ... }
hfchat:   { text_only: true, ... }      // BLOCKED from images

// image (line ~2095+)
{ id: 'huggingface', supportsImg2Img: true, supportsInpainting: true,
  supportsImageInput: true, supportsImageOutput: true, ... }
{ id: 'realesrgan', supportsUpscale: true, taskOnly: true,
  supportsImageOutput: true, ... }      // upscale-only
```

Pre-call regex `HF_TEXT_ONLY_MODELS` blocks Mistral / Llama / Qwen / Gemma / Phi / Falcon / MPT / Flan / T5 / BART / Pythia / DeepSeek / Yi / Olmo from ever reaching `hfImageCall`.

6.2 withTimeout + AbortController

```
function withTimeout(promiseOrThunk, ms, label) {
  const ac = new AbortController();
  const p = (typeof promiseOrThunk === 'function')
    ? promiseOrThunk(ac.signal) : promiseOrThunk;
  const timer = new Promise( (_, reject) => {
    setTimeout(() => { ac.abort(); reject(...timed out...); }, ms);
  });
  return Promise.race([Promise.resolve(p).finally(clearTimeout), timer]);
}
```

All provider calls wrapped: 45 s LLM / 60 s fallback chain / 120 s image gen. Loading state always clears in `finally{}`.

6.3 Output Receipt audit log

Every successful image task records to localStorage `ps_receipts` (rolling 20):

```
{ ts, task, providerId, providerName, prompt, seed, consistencyMode,
  skipped: [...reasons], timedOut: false }
```

Visible in Settings → Output Receipts. Audit trail when results look wrong.

6.4 Multi-turn / Scene Lock

After every successful gen: `setCurrentImageFromUrl(resultUrl, 'Previous result – keep editing')` updates `img` + `lastSDPrompt`. Next turn's prompt gets injected with 1000-char preservation directive listing subject identity / clothing / pose / lighting / palette etc. when `sceneLocked` is true.

7. What is shipped (v17dn → v17e9)

Version	Commit	Headline change
v17dn	36a2387	Restore v17dg + 5 free providers (12 image total)

v17do	e7ba766	Pollinations Flux promoted to slot 1 (blur fix)
v17dp	a7fb19e	Scene Lock 1000-char preservation directive
v17dq	06afb66	Task router · character slot · sharpening · export-prompt fallback · Local SD slot
v17dr	cdfe88a	runImageTask facade · soft img2img → t2i fallback · smart cleanPrompt
v17ds	b5da6b0	ADD_OBJECT hard-route · MANUAL_MASK_NEEDED · similarity reject
v17dt	9d2e237	HF text-only regex block · supportsImageOutput / Input / Inpainting gates
v17du	a0e7381	withTimeout 45/60/120s · AbortController · skipped-reasons surfaced · debug line
v17dv	da4f5bd	ADD_OBJECT auto-placement bounding box (sky / ground / on-subject / default)
v17dw	7ee8ee0	Test Keys diagnostic (masked + status icons)
v17dx	1d05375	Output Receipts audit log
v17dy	a1e1a09	Test Keys grouped image vs chat-only · TEXT-ONLY locked labels
v17dz	6fe7f0d	HF Mistral hidden from fallback display when image attached
v17e0	593d410	Ship #1 Manual Mask painter (USER VERIFIED — sunglasses test)
v17e1-e3	(reverted)	FAILED — bg-removal crashed iPad twice; sentinel-throw broke load
v17e4-e5	2fff569 / ce42ed8	Revert to v17e0 source byte-identical · cache forward
v17e6	2027cc3	SiliconFlow + Real-ESRGAN + GFPGAN + CodeFormer providers (untested)
v17e7	04febd7	reply.match string coercion fix
v17e8	1bd6cea	callPuter response normalized at source
v17e9	c210b39	Edit regression fix · taskOnly flag respected · supportsFaceRestore added

8. What is verified vs not

USER-VERIFIED: v17e0 only. Sunglasses test (paint mask + "add sunglasses" → result with un-painted area preserved).

SHIPPED but NOT user-verified: v17e1-v17e9 features. Site-loading was confirmed at v17e0 (post-revert), v17e8 (after reply.match fix), and current tip. Individual feature behavior of v17e6+ providers and v17e9 routing has NOT been tested by the user.

FAILED outright: background removal — both OpenCV.js (8 MB WASM crashed iPad) and HF briaai/RMBG-1.4 (also crashed). PARKED.

9. Open roadmap (37 items)

From `PLAN_LOAD_AI.md` Section "Merged forward roadmap".

#	Item	Source
1	Manual Mask painter SHIPPED v17e0	spec Ph 2
2	Background removal PARKED	OpenCV / HF rembg crashed
3	rembg via Transformers.js + WebGPU	spec Ph 2
4	More Jimp ops (resize / crop / flip / compose)	original K
5	Output Verification (vision diff + auto-retry)	spec Ph 1.E
5b	Identity Lock (IP-Adapter / InstantID via HF Spaces)	user 2026-04-30
5c	Structured prompt-rewrite layer (JSON output)	user 2026-04-30
6	Character Cards UI (multi-profile)	spec Ph 3
7	3-mode toggle UI (strict / moderate / loose)	spec Ph 3
7a / 7b / 7c	Add / Remove / Merge images explicit flows	user 2026-04-30 · Wave 3
8	Lottie loading animations	original K
9	Voice-driven editing (Web Speech)	original M
10-13	Verse → Scene · Image → Verse · Visualize as I read · One-tap Make a Book Cover	original L (Tier-4 biblical)

14a/b/c/d	Curated style library · Face restore wiring · Real-ESRGAN routing · SiliconFlow verification	Glam parity
15	HF Spaces connector (public Gradio APIs)	spec Ph 2
16	Cohesive icon set across ACR / Attain / Attain Jr / Study	original E
17	App Store readiness (NSFW / watermark / privacy / vault)	original N
18a-e	Image → Video — basic motion / SVD / AnimateDiff / Talking Avatar / Effects	user 2026-04-30 KEY FEATURE
18-fallback	External video-prompt fallback (Runway / Kling / Pika / Luma / Wan)	Glam doc
19a	Batch edits (Wave 2)	user 2026-04-30 · original Wave 2
20-26	Tier-4 backlog (timeline · multi-language · wallpaper · daily inspiration · auto alt-text · offline queue · style translation)	original L
27-35	Polish backlog (A/B compare · history gallery · cache · auto-enhance · prompt enhance · camera roll · share-target · recipe · health UI)	original Tier 2/3
OPT	ComfyUI / A1111 / Fooocus / InvokeAI on companion machine	opt-in only
OPT-track	Step1X-Edit · Qwen-Image-Edit · FLUX Kontext · HF SDXL Inpainting · HF SD-1.5 Inpainting	future-track when free endpoints exist

10. How to add a feature (concrete walkthrough)

1. Pick a single small unit — one file, ~50 lines max, no new dependencies, no WASM
2. Branch off `main` ; never edit unless `git status` is clean
3. Add the code with comments explaining WHY (not WHAT — the code shows that)
4. Bump cache strings **forward**: `load/sw.js` + `load/image-prompt/sw.js` + version badge in `load/load.js`
5. Run `node --check <extracted-script>.js` for syntax sanity
6. Commit with detailed body: what, why, where, recovery command
7. Push to `main`
8. **Wait for user verification on iPad.** Don't ship the next item until they confirm "verified" or "broken".

9. If verified: append a row to `VERIFIED_LOG.md` (version, SHA, browser, exact test, what passed, what NOT tested), create `backup/<date>-v<ver>` branch.
10. If broken: **surgical revert only** — see Section 12.

11. How to add a provider

Locate `IMAGE_PROVIDERS = [` in `load/image-prompt/index.html` (~line 2095).
Add a new entry following the established shape:

```
{
  id: 'newprovider',
  name: 'New Provider',
  needsKey: true,
  enabled: () => !!C.newProviderKey && C.useProvider.newprovider !== false,
  supportsImg2Img: true,      // if true, edits route here
  supportsImageInput: true,
  supportsImageOutput: true,
  supportsInpainting: false,  // mask-aware?
  supportsUpscale: false,
  supportsFaceRestore: false,
  taskOnly: false,           // true = only run when filterImageProvidersForTask matches a specialty
  gen: async (prompt, imgData, opts) => {
    opts = opts || {};
    const r = await fetch('https://api.example.com/v1/generate', {
      method: 'POST',
      headers: { 'Authorization': 'Bearer ' + C.newProviderKey, 'Content-Type': 'application/json' },
      body: JSON.stringify({ prompt: prompt.slice(0, 2000), ...opts.mask && { mask: opts.mask.base64 } })
    });
    if (!r.ok) throw new Error(await readError(r, 'New Provider'));
    const data = await r.json();
    if (!data.image_url) throw new Error('New Provider empty response');
    const imgR = await fetch(data.image_url);
    const blob = await imgR.blob();
    if (!blob.type.startsWith('image/')) throw new Error('New Provider non-image');
    return URL.createObjectURL(blob);
  }
}
```

Then add the key field + toggle to the Settings panel HTML; add the corresponding entries to `useProvider` defaults (line ~620), `saveSettings` (line ~1070), and `loadSettings` (line ~960). Test Keys diagnostic auto-pulls if `checkOneKey` recognises the provider id (line ~620).

12. Surgical revert procedure

From CLAUDE.md (mandatory rule). When something breaks, NEVER blanket-reset.

1. `git diff <good-sha>..<bad-sha> -- <file>` — see exactly what changed
2. `git checkout <good-sha> -- <specific-file>` — restore ONLY the broken file
3. For shared files: hand-port the un-broken hunks forward; reverse only the breaking ones
4. Before touching `main`: list every `VERIFIED_LOG.md` row between target and HEAD; each must survive the revert or be re-applied
5. Cache strings always go FORWARD, never backward (going back leaves iOS Safari with the broken old SW running)
6. Document the revert in `VERIFIED_LOG.md`: what reverted, what preserved, why surgical, recovery commands
7. If full-tree revert is genuinely needed: `git branch pre-revert-<date>-v<ver>` on current HEAD FIRST, then revert

13. Cache discipline + deploy flow

Every JS / HTML / CSS edit MUST bump:

- `load/sw.js` — `var CACHE = 'load-vXXY'` (alpha-incremented forward)
- `load/load.js` — the on-screen badge `<span id="ve-version">vXXY</span>`
- `load/image-prompt/sw.js` — `const CACHE = 'image-prompt-vNN'`
- Same pattern for ACR (`acr-vNN`), Attain (`attain-vNN`), Attain Jr (`attain-jr-vNN`), Study (`acr-study-vNN`)

Skipping the bump = users serve stale cached code. Going BACKWARDS in the cache name = iOS Safari may keep the old broken SW alive (don't).

Deploy: push to `main` on GitHub. Pages picks it up; CDN edge cache lags 5-15 min. Always cache-bust verify URLs: `?bust=<version>`.

14. Backup branches

30 backup branches on `origin` . Each captures the **entire repo** (every site at that moment).

Branch	Purpose
<code>backup/2026-04-30-v17e0</code>	Last user-verified working state (sunglasses test)
<code>backup/2026-04-30-v17dn</code>	Stable 12-provider chain image-gen
<code>backup/2026-04-29-v17dg</code>	Pre-spec verified baseline (7 truly-free providers)
<code>backup/main-2026-04-28</code>	Snapshot of <code>main</code> before this session's spec build wave
<code>backup/2026-04-29-acr-v17</code> · <code>2026-04-28-acr-v16</code>	ACR Reader milestones

Recover via `git checkout backup/<name>` . Forward-only restore to `main` :

```
git checkout main
git checkout backup/2026-04-30-v17e0 -- .
git commit -m "restore: roll main to v17e0"
git push origin main
```

15. Source-of-truth doc map

File	Purpose	Update when
<code>PLAN_LOAD_AI.md</code>	Merged build plan (provider lock list · routing · 37-item roadmap · Editor Tools section · wave map)	End of every shipped version
<code>VERIFIED_LOG.md</code>	Per-version verification ledger	The MOMENT user signs off, never end-of-day
<code>HANDOFF.md</code>	Long-running cross-session task list (locked deliverable commitments for Attain / Attain Jr extractions, free hosting decisions)	Major architectural shifts
<code>CLAUDE.md</code>		When user adds a new rule

	Rules for AI / human devs (session continuity · backups · cache discipline · surgical revert · no-guessing · sandbox gaps)	
<code>SESSION_NOTES_YYYY-MM-DD.md</code>	Per-day journal — current state · built today · outstanding · today's commits	End of every session
<code>inbox/PLAN_LOAD_AI.md</code>	Mirror of build plan for user's iPad GitHub view	Whenever main copy changes
<code>inbox/HANDOFF_README.md</code>	60-second briefing for next dev	End of session
<code>inbox/*.docx</code> · <code>*.pdf</code> · <code>*.zip</code>	User-uploaded source materials (Anti-Blur · Lock Providers · Fix Hang · Key Diagnostic · Glam Style Research · character-system zip)	Don't delete — historical reference

16. Gotchas / capability gaps

- **iPad Safari + WASM:** 8 MB+ WASM bundles can crash the tab. OpenCV.js is OUT. Transformers.js + WebGPU might work but needs strict memory testing first.
- **iOS Safari SW eviction:** if cache name doesn't change (or goes BACKWARD), the broken old SW stays active. ALWAYS bump cache forward.
- **GitHub Pages CDN lag:** 5-15 min from push to live. Use `?bust=<ver>` for verification.
- **Sandbox host blocks (varies per session):** `dssorit.github.io` often blocked from this build server. Use `raw.githubusercontent.com/...` or GitHub MCP `get_file_contents` for verification.
- **Puter SDK response shape:** can return string / `{message:{content}}` / `{message:{content:array}}` / `{content}` / `{text}` / array. Always normalize at `callPuter` source (line ~1751).
- **HF cold-start:** first call to a model may return 503 with "model loading". `hfImageCall` throws "Model loading (~30s cold start)" — caller must handle.
- **Anthropic in browser:** needs `anthropic-dangerous-direct-browser-access` flag which exposes the key. Excluded from default chat chain; opt-in only.
- **VPN traffic:** some VPNs throttle GitHub raw downloads. The user has VPN active in screenshots — large downloads may need VPN off.
- **Free-tier providers vary:** SiliconFlow / Together / Replicate / DeepInfra / Fal — credits run out without warning. Mark opt-in only.

- **Real-ESRGAN / GFPGAN / CodeFormer on HF Inference:** some are Spaces-only, not Inference API. May 404. The taskOnly flag (v17e9) ensures they only run when explicitly routed.

17. Recommended next steps

1. **Verify v17e9 features end-to-end on iPad** — SiliconFlow opt-in test, "smooth my skin" → GFPGAN routing, "upscale 4x" → Real-ESRGAN. If any fail, surgical revert just that provider's entry, leave the rest.
2. **Item #4 — More Jimp ops (resize / crop / flip / compose).** Pure JS, ~150 KB, already loaded. Safest next ship. ~30 lines added to `LOCAL_OPS` array.
3. **Item #5 — Output Verification (Pipeline E).** Closes the spec's last unimplemented pipeline. After every gen, vision LLM compares result vs source, returns identity / style / instruction match scores. Auto-retry once if below threshold. Hooks into existing `imageSimilarity`.
4. **Item #18a — Image → Video tier 1 (basic motion).** Browser-only Ken-Burns pan/zoom/parallax via Canvas + Anime.js + FFmpeg.wasm. Zero AI cost; first concrete video output.
5. **Items #6-7 — Character Cards + 3-mode toggle UI.** Productizes the existing CHAR slot; spec-required.
6. **Items #10-13 — Tier-4 biblical features.** Cross-app integration with ACR Reader; the project's unique differentiator.

Decision rule for next ship: smaller surface area first, biggest blast radius items last. The user has been burned by big changes; rebuild trust with verified small wins.